






Does more calibration data help?

Exploratory reuse of five previously evaluated artist splits



Below zero: better than raw probabilities

Lines: median across calibration subsets
Shading: 10th-90th percentile subset spread
(not confidence intervals)

24 / 47 / 71 artists: 30 nested repeats
94 artists: one shared full-pool fit

Models and evaluation tracks stay fixed
within each split.